

Latency of ML inference in hot environment
NVIDIA A100, n=100, 95% confidence

