

Runtime of ML inference in hot environment  
NVIDIA A100, n=100, 95% confidence

