

Latency of minimal CUDA API execution  
NVIDIA A100, n=100, 95% confidence

